

Sample Efficient Prioritized Experience Replay for Adversarial Environments

Tianyang Duan, Zongyuan Zhang, Zekai Sun, Heming Cui

Motivation

➤ Adversarial environment

- Non-stationary environment
- Zero-sum game
- Rapidly changing conditions and unpredictable opponents

➤ Reinforcement learning(RL) shows promising results in adversarial environments.


- Learn through interactions between agents and environments.
- Require a large number of trials and errors to converge

➤ Sample efficiency of RL needs to be improved in adversarial environments.

- Reduce the amount of data required for reaching target performance level.



Robot soccer



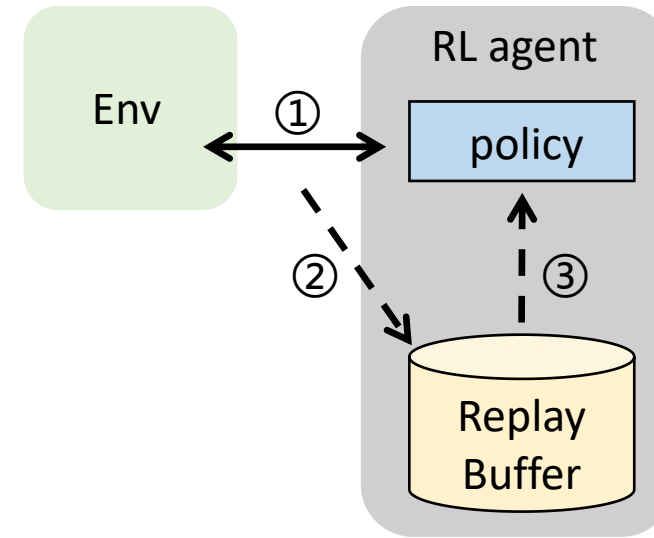
RoboSumo

Experience Replay


➤ **Experience replay** stores past experiences in a replay buffer and sampled randomly during training to improve the sample efficiency.

- ① Agent interacts with the environment.
 - Agent receives a state S and out put action A
 - Environment transforms to state S' and give a reward R
- ② Store the transaction (S, A, R, S') in replay buffer.
- ③ Sample a batch of transactions to train.

- **However, random sampling degrade sample efficiency, especially in non-stationary environments.**



Existing Prioritized Experience Replay Method



➤ **Prioritized Experience Replay (PER)** gives priority to the "important" transactions.

- $Q(S_t, A_t)$: state-action value function, denotes the expected return with state S_t and action A_t

$$Q(S_t, A_t) \approx \sum_{k=t}^T R_k = R_t + \sum_{k=t+1}^T R_k \approx R_t + Q(S_{t+1}, A_{t+1})$$

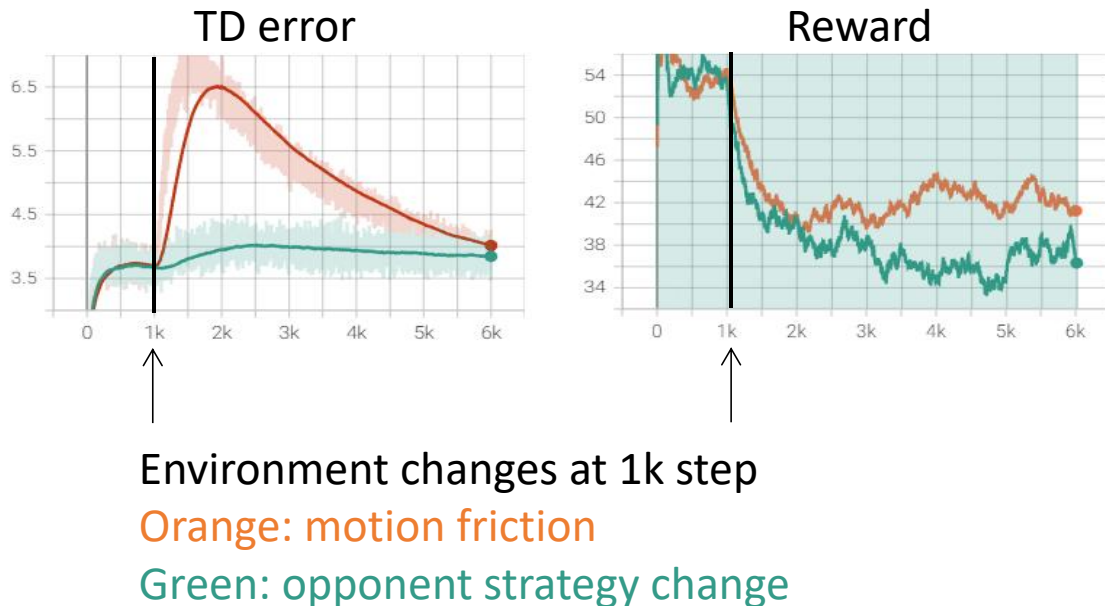
- Temporal Difference (TD) : The difference of expected and actual state-action value

$$TD_error = R_t + Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)$$

- Priority is calculated based on the TD error of the transaction.
- Why TD error: TD error reflects the degree of deviation of a sample from the policy

Problems of Existing PER Method in Adversarial Environment

High TD error ← State-action space edge
← Changes in environment



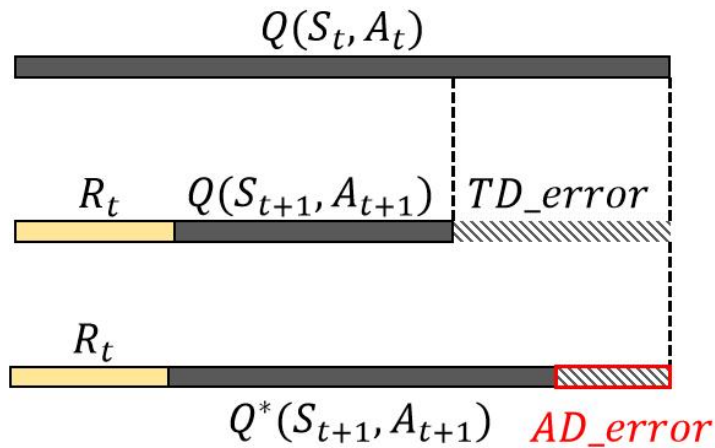
- Prioritized transactions at the edge of the state-action space (with high TD error) hinder the fast convergence of RL[1].
- In the later training stages, outdated samples are prioritized which introduce noise.
- TD error cannot fully capture opponent strategy changes.
- **Thus, TD-error based PER approach is not ideal for adversarial environments**

From Temporal Difference to Adversarial Difference

- We introduce the concept of Adversarial Difference (AD) as a measure to compare the expected returns from old and new opponents.
- Why AD:
 - TD indicates deviation of a transaction **versus the same opponent**
 - AD indicates deviation of a transaction **versus the different opponent**
- We propose **Adversarial Prioritized Experience Replay (APRE)** based on AD.
 - For improving sample efficiency in adversarial environments
 - APRE faces the following challenges:
 - ① How to define old and new transactions
 - ② How to trade off between old and new transaction
 - new experience: limited amount, help quickly converge
 - old experience: improving stability, avoiding catastrophic forgetting and overfitting

Key Idea: Adversarial Difference

- We propose Adversarial Difference(AD) which compares the expected returns between the old opponent and the new opponent.
- A transaction $(S_t, A_t, R_t, S_{t+1}, A_{t+1})$
 - $AD_error = R_t + Q^*(S_{t+1}, A_{t+1}) - Q(S_t, A_t)$
 - Q : the state-action value function of current opponent, Q^* : the state-action value function of old opponent



$$Q(S_t, A_t) \approx \sum_{k=t}^T R_k = R_t + \sum_{k=t+1}^T R_k \quad \leftarrow \text{Expected return versus new opponent}$$

$$TD_error = R_t + Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \quad \leftarrow \text{The difference of expected return at time step } t \text{ and } t + 1 \text{ versus new opponent}$$

$$Q^*(S_{t+1}, A_{t+1}) = \sum_{k=t+1}^T R_k^* \quad \leftarrow \text{Expected return versus old opponent}$$

$$AD_error = R_t + Q^*(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \quad \leftarrow \text{The difference of expected return versus old opponent and new opponent}$$

$$\approx \sum_{k=t+1}^T R_k^* - \sum_{k=t+1}^T R_k$$

APER: Adversarial Prioritized Experience Replay

➤ **Ratio sampling** from “new” section and “old” section

- Stage1: Start by detecting a rise in TD error, λ start from 1 and decays with training steps
- Stage2: Starting from $\text{new}/\text{sum} > \lambda$
- Stage3: *Starting from* $\text{new}/\text{sum} = 1$

➤ **Prioritized sampling** from each individual section

In Stage 1,2:

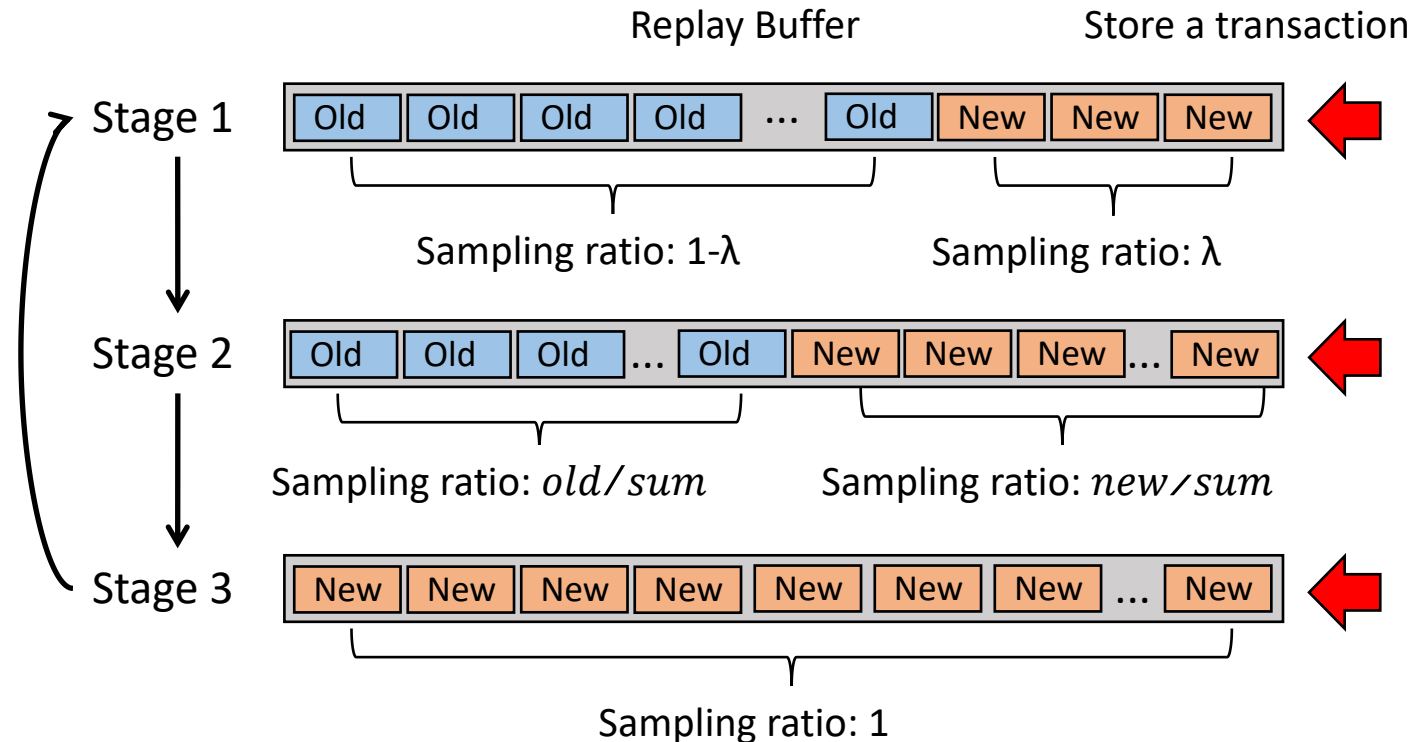
- old sample priority $\propto 1/|\text{AD_error}|$
- new sample priority $\propto |\text{AD_error}|$

Stage 3:

- priority $\propto |\text{TD_error}|$



The rise of TD error can be observed when the opponent change



Implementation and Evaluation

➤ Evaluation environment:

- rSoccer - IEEE VSSS environment to simulate a robot soccer scenario.

➤ RL algorithm:

- TD3 (Twin Delayed DDPG): a state-of-the-art RL algorithm for continuous control tasks

➤ Baseline:

- Vanilla-ER
- Prioritized Experience Replay (PER)
- Combined Experience Replay (CER)

➤ Evaluation settings:

- 3 vs 3 robot soccer match
- Each agent controls one player
- Follow the rewards shaping in rSoccer
- Each match lasting 200 time steps



Evaluation Questions


- Does the method improve sample efficiency in adversarial environments?
- How does AD-error-based priority help to reduce the efficiency damage caused by outdated transactions?
- What is the method's sensitivity to critical sampling-related hyperparameters?

End-to-end Performance


- **APER** (Red line) take a minimum of episodes to restore the average reward before changing the opponent policy, with a **46.64% improvement to Vanilla ER**.
- Play in the above experimental environment
- At 5000 episodes, the opponent's strategy changed from offensive to defensive
- All methods start from the same pre-trained model
- All methods use the same hyperparameters, with learning rate = 1×10^{-4} , batch size = 1024, buffer size = 5×10^5 .

Replay	Episodes spent to restore the original rewards	Improvement over vanilla ER
APER (Ours)	14695	46.64%
CER	19698	28.48%
PER	22064	19.89%
Vanilla ER	27543	-

Reward



Goal



Conclusion

- In this talk, we present APER, a sample efficient prioritized experience replay for adversarial environments.
 - We propose adversarial difference (AD), which compares the expected gains between different opponents.
 - APER detects the changes of the opponent by the rise of TD error.
 - APER uses ratio sampling to trade off between the new transactions and old transactions.
 - APER samples from new and old transactions according to AD and $1/AD$ priority, respectively.
- APER requires a minimum of episodes to restore the average reward before changing the opponent policy compared to baselines.
- Plan to submitted to 26th European Conference on Artificial Intelligence (ECAI 23) in April.

Future work

➤ Short-term work

- Evaluate APER in other adversarial environments (e.g. RoboSumo, Pong)
- Evaluate APER with other RL algorithms (e.g. A2C,DDPG)

➤ Long-term work

- Apply APER in real robotic multi-agent reinforcement learning environments
 - Experience replay sharing in distributed training
- Apply APER in Deep Reinforcement Learning